

## Problem Statement

Ninjacart is India's largest fresh produce supply chain company. They are pioneers in solving one of the toughest supply chain problems of the world by leveraging innovative technology. They source fresh produce from farmers and deliver them to businesses within 12 hours. An integral component of their automation process is the development of robust classifiers which can distinguish between images of different types of vegetables, while also correctly labeling images that do not contain any one type of vegetable as noise.

As a starting point, ninjacart has provided us with a dataset scraped from the web which contains train and test folders, each having 4 sub-folders with images of onions, potatoes, tomatoes and some market scenes. We have been tasked with preparing a multiclass classifier for identifying these vegetables. The dataset provided has all the required images to achieve the task.

## Importing the dataset and doing usual exploratory analysis steps like checking the structure & characteristics of the data (10 points)

Visualize the data, use the dataset directory to create a list containing all the image paths in the training folder. You can use matplotlib or tensorflow to plot a grid sample of the images you fetched from the list of image paths.

Plot a few of the images of each class to check their dimensions.  
[Note that the images are not all of uniform dimensions]

Verify the count of images in each train and test folder by plotting histogram .

Check each folder to see if the number of images matches the reported number.

## Testing your best model so far(20 points)

Testing on the test set & Random image samples prediction[10]

Summary & Insights [10]

```
!gdown 1clZX-lV_MLxKHSyeyTheX50CQtNCUcqT  
#!/unzip /content/ninjacart_data.zip -d /content/New_Folder/1
```

Downloading...

From (original): [https://drive.google.com/uc?id=1clZX-lV\\_MLxKHSyeyTheX50CQtNCUcqT](https://drive.google.com/uc?id=1clZX-lV_MLxKHSyeyTheX50CQtNCUcqT)

From (redirected): [https://drive.google.com/uc?id=1clZX-lV\\_MLxKHSyeyTheX50CQtNCUcqT&confirm=t&uuid=f2d72ee4-94c4-437f-8222-](https://drive.google.com/uc?id=1clZX-lV_MLxKHSyeyTheX50CQtNCUcqT&confirm=t&uuid=f2d72ee4-94c4-437f-8222-)

```
a2bc5a773b2c
To: /content/ninjacart_data.zip
100% 275M/275M [00:02<00:00, 99.7MB/s]

#import shutil
#shutil.rmtree('ninjacart_data/')

!unzip -q ninjacart_data.zip

#!pip install tensorflow

# Tensorflow import
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.layers import Conv2D, MaxPooling2D,
GlobalAveragePooling2D, Dense, ReLU, Softmax, BatchNormalization,
Dropout
from tensorflow.random import set_seed

import os

import matplotlib.pyplot as plt
import pandas as pd
import random
import glob
import sklearn.metrics as metrics

class_dirs = os.listdir("/content/ninjacart_data/train") # list all
directories inside "train" folder
image_dict = {} # dict to store image array(key) for every
class(value)
count_dict = {} # dict to store count of files(key) for every
class(value)
count_dict_test={}
print('Total Classes : ',class_dirs)
# iterate over all class_dirs
for cls in class_dirs:
    # get list of all paths inside the subdirectory
    file_paths = glob.glob(f'/content/ninjacart_data/train/{cls}/*')
    #print(file_paths)
    count_dict[cls]=len(file_paths)
    image_path=random.choice(file_paths)
    image_dict[cls]=tf.keras.utils.load_img(image_path)

print('Training - Total Images : \n',count_dict)
#print('Test - Total Images : ',len(test_ds))

#print(image_dict.keys())
plt.figure(figsize=(15,8))
```

```

for i,(cls,img) in enumerate(image_dict.items()):
    plt.subplot(3,2,i+1)
    plt.imshow(img)
    plt.axis('off')
    plt.title(f'{cls},{img.size}')
for cls in class_dirs:
    # get list of all paths inside the subdirectory
    file_paths = glob.glob(f'/content/ninjacart_data/test/{cls}/*')
    count_dict_test[cls]=len(file_paths)

print('Test - Total Images : \n',count_dict_test)

```

```

Total Classes : ['onion', 'tomato', 'indian market', 'potato']
Training - Total Images :
{'onion': 849, 'tomato': 789, 'indian market': 599, 'potato': 898}
Test - Total Images :
{'onion': 83, 'tomato': 106, 'indian market': 81, 'potato': 81}

```

onion,(183, 275)



tomato,(400, 500)



indian market,(275, 183)



potato,(300, 168)



## Exploratory Data Analysis. (20 points)

Plotting class distribution & Visualizing Image dimensions with their plots[10]

Splitting the dataset into train, validation, and test set[10]


```

plt.figure(figsize=(10,3))
plt.subplot(121)
plt.bar(count_dict.keys(),count_dict.values(),label=count_dict.values(
))
plt.title('Training Dataset')
plt.subplot(122)

```

```
plt.bar(count_dict_test.keys(),count_dict_test.values())
plt.title('Test Dataset')

Text(0.5, 1.0, 'Test Dataset')
```



Split the dataset to a train and validation set.

The provided data does not contain separate training and validation folders. For us to do hyperparameter tuning of our models, it is important to divide the dataset into an 80-20 split for training and validation respectively.

```
#pip install split-folders
```

```
'''import splitfolders
try:
    splitfolders.ratio("/content/ninjacart_data/train",seed=1337,
output="ninjacart_data/Train", ratio=(0.8, 0.2))
except:
    pass'''
```

```
{"type": "string"}
```

Before fitting data to our model, we must make sure that each image is square-shaped so that we may resize it to the required dimensions and also perform rescaling which will rescale the inputs between 0-1 by dividing each value by 255.

```
image_size = (128, 128)
train_ds,val_ds=tf.keras.utils.image_dataset_from_directory('/content/
ninjacart_data/train',label_mode = 'categorical',image_size =
image_size,validation_split=0.2,subset='both',shuffle=True,seed=42)
test_ds=tf.keras.utils.image_dataset_from_directory('/content/ninjacar
t_data/test',label_mode = 'categorical',image_size = image_size)
```

```

#val_ds=tf.keras.utils.image_dataset_from_directory('/content/ninjacar
t_data/train',image_size =
image_size,validation_split=0.2,subset='validation',shuffle=True,seed=
42)

height,width=128,128

Found 3135 files belonging to 4 classes.
Using 2508 files for training.
Using 627 files for validation.
Found 351 files belonging to 4 classes.

'''train_ds=tf.keras.utils.image_dataset_from_directory('/content/
ninjacart_data/Train/train')
val_ds=tf.keras.utils.image_dataset_from_directory('/content/ninjacart
_data/Train/val')
test_ds=tf.keras.utils.image_dataset_from_directory('/content/ninjacar
t_data/test')

height,width=128,128
data_preprocess_with_flatten=keras.Sequential(
    name='data_preprocess_with_flatten'
    ,layers= [
        layers.Resizing(height,width),
        layers.Rescaling(1/255),
        #layers.Flatten(),
    ]
)

# Perform Data Processing on the train, val, test dataset
train_ds = train_ds.map(lambda x, y: (data_preprocess_with_flatten(x),
y))
val_ds = val_ds.map(lambda x, y: (data_preprocess_with_flatten(x), y))
test_ds=test_ds.map(lambda x,y :
(data_preprocess_with_flatten(x),y))'''

{"type": "string"}

sample = next(iter(train_ds))[0]

#print(sample)

```

## Creating model architecture and training (50 points)

Defining the CNN Classifier model from scratch[10]

Improving Baseline CNN to reduce overfitting[10]

Implementing Callbacks while training the model[10]

Finetune pretrained models such as VGG, ResNet and MobileNet[10]

Plotting the model training metrics and confusion matrix[10]

```
#from tensorflow.keras import layers
#from tensorflow.keras import regularizers

def plot_accuracy(model_fit):
    #accuracy graph

    x = range(0, len(model_fit.history.history['accuracy']))
    y_train = [acc * 100 for acc in
model_fit.history.history['accuracy']]
    y_val = [acc * 100 for acc in
model_fit.history.history['val_accuracy']]

    plt.plot(x, y_train, label='Train', color='b')
    #annot_max(x, y_train, xytext=(0.7,0.9))
    plt.plot(x, y_val, label='Val', color='r')
    #annot_max(x, y_val, xytext=(0.8,0.7))
    plt.ylabel('Accuracy', fontsize=15)
    plt.xlabel('epoch', fontsize=15)
    plt.legend()
    plt.show()

import seaborn as sns
def ConfusionMatrix(model, ds, label_list):
    # Note: This logic doesn't work with shuffled datasets

    # generate confusion matrix and plot it
    cm =
metrics.confusion_matrix(true_categories, predicted_categories) # last
batch
    sns.heatmap(cm, annot=True, xticklabels=label_list,
yticklabels=label_list, cmap="YlGnBu", fmt='g')
    plt.show()

noise_path = '/content/ninjacart_data/test/indian market'
onion_path = '/content/ninjacart_data/test/onion'
potato_path = '/content/ninjacart_data/test/potato'
tomato_path = '/content/ninjacart_data/test/tomato'
def classwise_accuracy(class_path, class_name, model_name) :
    paths = []
    for i in os.listdir(class_path):
        paths.append(class_path + "/" + str(i))

    correct = 0
    total = 0

    for i in range(len(paths)):
        total+= 1
```

```

img = tf.keras.utils.load_img(paths[i])
img = tf.keras.utils.img_to_array(img)
img = tf.image.resize(img, (128, 128))
img = tf.expand_dims(img, axis = 0)

pred = model_name.predict(img, verbose=0)
if tf.argmax(pred[0]) == class_names.index(f"{class_name}"):
    correct+= 1

print(f"Accuracy for class {class_name} is
{round((correct/total)*100, 2)}% consisting of {len(paths)} images")

def model_evaluation_acc(model_name):
    # Evaluate the model
    loss, acc = model_name.evaluate(test_ds, verbose=2)
    print("Restored model, accuracy: {:.2f}%".format(100 * acc))

    y_pred = model_name.predict(test_ds)
    predicted_categories = tf.argmax(y_pred, axis=1)
    true_cat = tf.concat([y for x, y in test_ds], axis=0)
    true_categories = tf.argmax(true_cat, axis=1)

    # calculate accuracy
    test_acc = metrics.accuracy_score(true_categories,
predicted_categories) * 100
    print(f'\nTest Accuracy: {test_acc:.2f}%\n')

    classwise_accuracy(noise_path, 'indian market', model_name)
    classwise_accuracy(onion_path, 'onion', model_name)
    classwise_accuracy(potato_path, 'potato', model_name)
    classwise_accuracy(tomato_path, 'tomato', model_name)

    ConfusionMatrix(model_name, test_ds, class_dirs)

cnn_model = tf.keras.Sequential(
    name='cnn_model',
    layers=[
        layers.Rescaling(1./255),
        layers.InputLayer(input_shape = [128, 128, 3]),

        layers.Conv2D(filters = 32, kernel_size = (3,3), padding = 'Same',
activation = 'relu'),
        layers.Conv2D(filters = 32, kernel_size = (3,3), padding = 'Same',
activation = 'relu'),
        layers.MaxPool2D(pool_size = (2,2)),

        layers.Conv2D(filters = 64, kernel_size = (3,3), padding = 'Same',
activation = 'relu'),
        layers.Conv2D(filters = 64, kernel_size = (3,3), padding = 'Same',
activation = 'relu'),

```

```

layers.MaxPool2D(pool_size = (2,2)),

layers.Conv2D(filters = 128, kernel_size = (3,3), padding =
'Same', activation = 'relu'),


layers.GlobalAveragePooling2D(),
layers.Dense(4, activation = 'softmax')
])
cnn_model.compile(optimizer='Adam',loss='categorical_crossentropy',met
rics='accuracy')
history=cnn_model.fit(train_ds,epochs=5,validation_data=val_ds)

Epoch 1/5
79/79 [=====] - 18s 128ms/step - loss: 1.0201
- accuracy: 0.5443 - val_loss: 0.7741 - val_accuracy: 0.6683
Epoch 2/5
79/79 [=====] - 12s 147ms/step - loss: 0.6591
- accuracy: 0.7552 - val_loss: 0.6689 - val_accuracy: 0.7400
Epoch 3/5
79/79 [=====] - 9s 111ms/step - loss: 0.6315
- accuracy: 0.7576 - val_loss: 0.5752 - val_accuracy: 0.7656
Epoch 4/5
79/79 [=====] - 9s 107ms/step - loss: 0.5820
- accuracy: 0.7775 - val_loss: 0.5651 - val_accuracy: 0.7544
Epoch 5/5
79/79 [=====] - 8s 90ms/step - loss: 0.5452 -
accuracy: 0.7863 - val_loss: 0.5467 - val_accuracy: 0.7608

plot_accuracy(cnn_model)

```





```
# Evaluate the model
class_names = ['indian market', 'onion', 'potato', 'tomato']
loss, acc = cnn_model.evaluate(test_ds, verbose=2)
print("Restored model, accuracy: {:.2f}%".format(100 * acc))

y_pred = cnn_model.predict(test_ds)
predicted_categories = tf.argmax(y_pred, axis=1)
true_cat = tf.concat([y for x, y in test_ds], axis=0)
true_categories = tf.argmax(true_cat, axis=1)

# calculate accuracy
test_acc = metrics.accuracy_score(true_categories,
predicted_categories) * 100
print(f'\nTest Accuracy: {test_acc:.2f}%\n')

classwise_accuracy(noise_path, 'indian market', cnn_model)
classwise_accuracy(onion_path, 'onion', cnn_model)
classwise_accuracy(potato_path, 'potato', cnn_model)
classwise_accuracy(tomato_path, 'tomato', cnn_model)

ConfusionMatrix(cnn_model, test_ds, class_dirs)

11/11 - 2s - loss: 0.6833 - accuracy: 0.6952 - 2s/epoch - 161ms/step
Restored model, accuracy: 69.52%
```

11/11 [=====] - 2s 38ms/step

Test Accuracy: 27.35%

Accuracy for class indian market is 43.21% consisting of 81 images

Accuracy for class onion is 78.31% consisting of 83 images

Accuracy for class potato is 49.38% consisting of 81 images

Accuracy for class tomato is 96.23% consisting of 106 images



```
test_dir = '/content/ninjacart_data/test'
test_images = []
class_names = ['indian market', 'onion', 'potato', 'tomato']

for folder in os.listdir(test_dir):
    for image in os.listdir(test_dir + '/' + folder):
        test_images.append(os.path.join(test_dir, folder, image))
img_0 = tf.keras.utils.load_img(random.choice(test_images))
img_0 = tf.keras.utils.img_to_array(img_0)
img_0 = tf.image.resize(img_0, (128, 128))
img_1 = tf.expand_dims(img_0, axis = 0)

def grid_test_model(model_name):
    fig = plt.figure(1, figsize=(8, 8))
    plt.axis('off')
```

```

n = 0
for i in range(4):
    n += 1

    img_0 = tf.keras.utils.load_img(random.choice(test_images))
    img_0 = tf.keras.utils.img_to_array(img_0)
    img_0 = tf.image.resize(img_0, (128, 128))
    img_1 = tf.expand_dims(img_0, axis = 0)

    pred = model_name.predict(img_1)
    predicted_label = tf.argmax(pred, 1).numpy().item()

    for item in pred :
        item = tf.round((item*100))

    plt.subplot(2, 4, n)
    plt.axis('off')
    plt.title(f'prediction : {class_names[predicted_label]}\n\n'
              f'{item[0]} % {class_names[0]}\n'
              f'{item[1]} % {class_names[1]}\n'
              f'{item[2]} % {class_names[2]}\n'
              f'{item[3]} % {class_names[3]}\n')
    plt.imshow(img_0/255)
    plt.show()
grid_test_model(cnn_model)

1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 17ms/step

<ipython-input-21-40cb159e3305>:32: MatplotlibDeprecationWarning:
Auto-removal of overlapping axes is deprecated since 3.6 and will be
removed two minor releases later; explicitly call ax.remove() as
needed.
    plt.subplot(2, 4, n)

1/1 [=====] - 0s 17ms/step

```

prediction : onion   prediction : potato   prediction : onion   prediction : onion

4.0 % indian market	1.0 % indian market	8.0 % indian market	1.0 % indian market
78.0 % onion	39.0 % onion	76.0 % onion	54.0 % onion
18.0 % potato	60.0 % potato	21.0 % potato	45.0 % potato
0.0 % tomato	0.0 % tomato	0.0 % tomato	0.0 % tomato



```

augmentation = tf.keras.Sequential([
    tf.keras.layers.RandomFlip("horizontal_and_vertical"),
    tf.keras.layers.RandomRotation(0.2),
    tf.keras.layers.RandomTranslation(height_factor = 0.2,
width_factor=0.2)
])

aug_ds = train_ds

for image, label in aug_ds :
    image = augmentation(image)

cnn_model_imp_performance=keras.Sequential(
    name='cnn_model_imp_performance',
    layers=[
        layers.Rescaling(1./255)

        ,layers.Conv2D(filters=32,kernel_size=3,padding='same',activation=
'relu',input_shape=(128,128,3))
        ,layers.BatchNormalization()
        ,layers.MaxPooling2D()

        ,layers.Conv2D(filters=32,kernel_size=3,padding='same',activation=
'relu')
        ,layers.BatchNormalization()
        ,layers.MaxPooling2D()

        ,layers.Conv2D(filters=64,kernel_size=3,padding='same',activation=
'relu')
        ,layers.BatchNormalization()
        ,layers.MaxPooling2D()

        ,layers.Conv2D(filters=64,kernel_size=3,padding='same',activation=

```

```

'relu')
    ,layers.BatchNormalization()
    ,layers.MaxPooling2D()

    ,layers.Conv2D(filters=128,kernel_size=3,padding='same',activation
='relu')
    ,layers.BatchNormalization()
    ,layers.GlobalAveragePooling2D()
#    ,layers.AveragePooling2D()
#    ,layers.Flatten()
#    ,layers.Dense(256,activation='relu')
    ,layers.Dropout(0.2)
    ,layers.Dense(4,activation='softmax')
])

```

```

cnn_model_imp_performance.compile(optimizer='Adam',loss='categorical_crossentropy',metrics='accuracy')
history=cnn_model_imp_performance.fit(train_ds,epochs=20,validation_data=val_ds)

```

Epoch 1/20

```

79/79 [=====] - 15s 125ms/step - loss: 0.6369
- accuracy: 0.7699 - val_loss: 1.2843 - val_accuracy: 0.4386

```

Epoch 2/20

```

79/79 [=====] - 8s 93ms/step - loss: 0.4591 -
accuracy: 0.8218 - val_loss: 1.4006 - val_accuracy: 0.2967

```

Epoch 3/20

```

79/79 [=====] - 7s 85ms/step - loss: 0.3978 -
accuracy: 0.8481 - val_loss: 1.2212 - val_accuracy: 0.4833

```

Epoch 4/20

```

79/79 [=====] - 9s 104ms/step - loss: 0.3658
- accuracy: 0.8557 - val_loss: 1.8263 - val_accuracy: 0.3573

```

Epoch 5/20

```

79/79 [=====] - 8s 99ms/step - loss: 0.3426 -
accuracy: 0.8636 - val_loss: 1.0784 - val_accuracy: 0.5933

```

Epoch 6/20

```

79/79 [=====] - 7s 85ms/step - loss: 0.3233 -
accuracy: 0.8792 - val_loss: 0.6574 - val_accuracy: 0.7416

```

Epoch 7/20

```

79/79 [=====] - 10s 118ms/step - loss: 0.3025
- accuracy: 0.8880 - val_loss: 2.5867 - val_accuracy: 0.3700

```

Epoch 8/20

```

79/79 [=====] - 10s 116ms/step - loss: 0.2858
- accuracy: 0.8915 - val_loss: 0.4279 - val_accuracy: 0.8102

```

Epoch 9/20

```

79/79 [=====] - 7s 85ms/step - loss: 0.2568 -
accuracy: 0.9067 - val_loss: 0.4054 - val_accuracy: 0.8501

```

Epoch 10/20

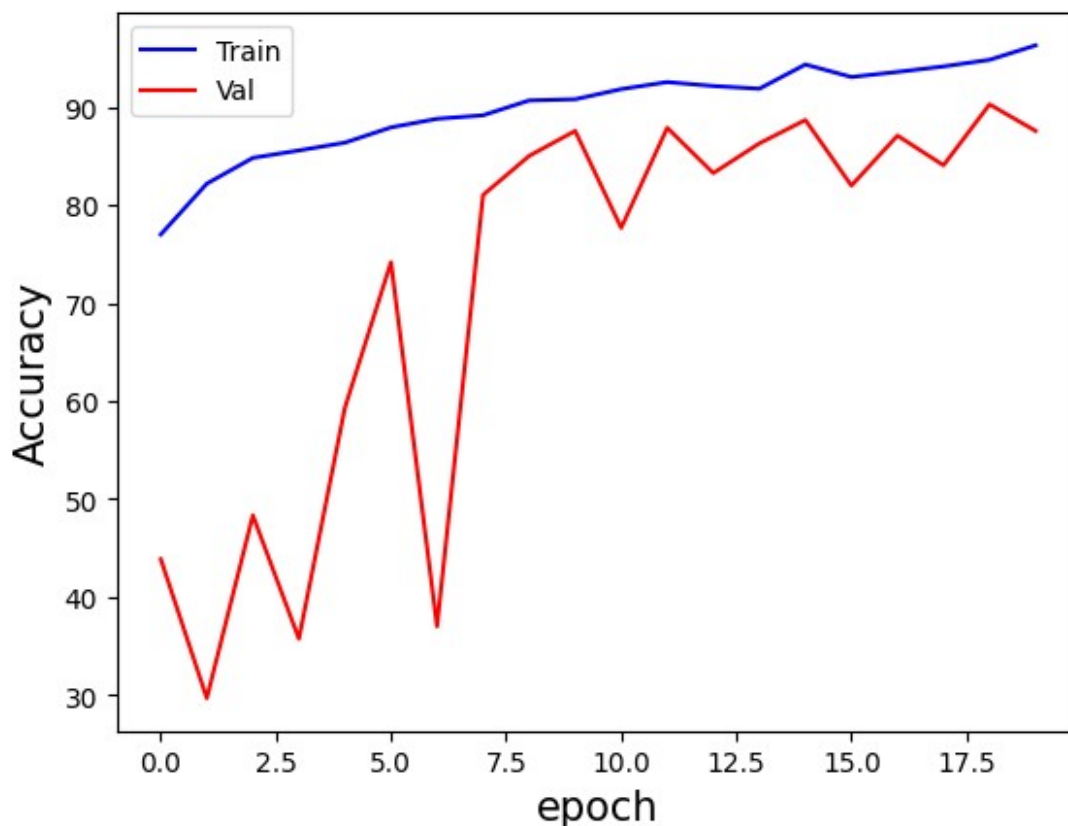
```

79/79 [=====] - 10s 114ms/step - loss: 0.2381
- accuracy: 0.9079 - val_loss: 0.2982 - val_accuracy: 0.8756

```

```
Epoch 11/20
79/79 [=====] - 9s 104ms/step - loss: 0.2266
- accuracy: 0.9183 - val_loss: 0.6611 - val_accuracy: 0.7767
Epoch 12/20
79/79 [=====] - 7s 85ms/step - loss: 0.2004 -
accuracy: 0.9254 - val_loss: 0.3064 - val_accuracy: 0.8788
Epoch 13/20
79/79 [=====] - 9s 102ms/step - loss: 0.2161
- accuracy: 0.9215 - val_loss: 0.4336 - val_accuracy: 0.8325
Epoch 14/20
79/79 [=====] - 9s 102ms/step - loss: 0.2114
- accuracy: 0.9187 - val_loss: 0.3485 - val_accuracy: 0.8628
Epoch 15/20
79/79 [=====] - 7s 84ms/step - loss: 0.1654 -
accuracy: 0.9434 - val_loss: 0.3499 - val_accuracy: 0.8868
Epoch 16/20
79/79 [=====] - 7s 85ms/step - loss: 0.1667 -
accuracy: 0.9306 - val_loss: 0.5411 - val_accuracy: 0.8198
Epoch 17/20
79/79 [=====] - 9s 103ms/step - loss: 0.1757
- accuracy: 0.9358 - val_loss: 0.3916 - val_accuracy: 0.8708
Epoch 18/20
79/79 [=====] - 8s 95ms/step - loss: 0.1587 -
accuracy: 0.9414 - val_loss: 0.5229 - val_accuracy: 0.8405
Epoch 19/20
79/79 [=====] - 7s 86ms/step - loss: 0.1393 -
accuracy: 0.9482 - val_loss: 0.2485 - val_accuracy: 0.9027
Epoch 20/20
79/79 [=====] - 9s 104ms/step - loss: 0.1122
- accuracy: 0.9629 - val_loss: 0.4119 - val_accuracy: 0.8756

plot_accuracy(cnn_model_imp_performance)
```



```
model_evaluation_acc(cnn_model_imp_performance)
```

```
11/11 - 1s - loss: 0.5283 - accuracy: 0.8063 - 1s/epoch - 106ms/step
```

```
Restored model, accuracy: 80.63%
```

```
11/11 [=====] - 1s 21ms/step
```

Test Accuracy: 24.50%

Accuracy for class indian market is 70.37% consisting of 81 images

Accuracy for class onion is 63.86% consisting of 83 images

Accuracy for class potato is 83.95% consisting of 81 images

Accuracy for class tomato is 98.11% consisting of 106 images



**We observe an accuracy jump of ~10 % by just :**

Applying augmentation to our data  
 Adding Dropout and BatchNormalization  
 Implementing callbacks during training

Use a model of your choice (could be vgg, resnet and mobilenet) and train it with an appropriate batch size.

Using the pretrained weights of popular networks is a great way to do transfer learning, since the size of our original dataset is small.

```
early_stopping_cb = tf.keras.callbacks.EarlyStopping(
    monitor = 'val_loss', patience = 5, restore_best_weights=True
)
pretrained_model=tf.keras.applications.VGG16(weights='imagenet',include_top=False,input_shape=(128,128,3))
pretrained_model.trainable=False
VGG16_model=tf.keras.Sequential([
    pretrained_model,
    #layers.Flatten(),
    layers.GlobalAveragePooling2D(),
```



```


        layers.Dense(4,activation='softmax')
    ])
VGG16_model.compile(optimizer='Adam', metrics='accuracy',
loss='categorical_crossentropy')

history=VGG16_model.fit(train_ds,epochs=20,validation_data=val_ds,callbacks=[early_stopping_cb],verbose=0)
#pretrained_model.summary()

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58889256/58889256 [=====] - 4s 0us/step

plot_accuracy(VGG16_model)

```



```

model_evaluation_acc(VGG16_model)


11/11 - 5s - loss: 1.0675 - accuracy: 0.8632 - 5s/epoch - 431ms/step
Restored model, accuracy: 86.32%
11/11 [=====] - 1s 48ms/step

Test Accuracy: 23.65%

Accuracy for class indian market is 85.19% consisting of 81 images

```

Accuracy for class onion is 89.16% consisting of 83 images  
 Accuracy for class potato is 67.9% consisting of 81 images  
 Accuracy for class tomato is 99.06% consisting of 106 images




```
'''early_stopping_cb = tf.keras.callbacks.EarlyStopping(
    monitor = 'val_loss', patience = 5, restore_best_weights=True
)'''
pretrained_model=tf.keras.applications.VGG19(weights='imagenet',include_top=False,input_shape=(128,128,3))
pretrained_model.trainable=False
VGG19_model=tf.keras.Sequential([
    pretrained_model,
    #layers.Flatten(),
    layers.GlobalAveragePooling2D(),
    layers.Dense(4,activation='softmax')
])
VGG19_model.compile(optimizer='Adam', metrics='accuracy', loss='categorical_crossentropy')

history=VGG19_model.fit(aug_ds,epochs=20,validation_data=val_ds,verbose=0)
#pretrained_model.summary()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5
80134624/80134624 [=====] - 5s 0us/step
```

```
plot_accuracy(VGG19_model)
model_evaluation_acc(VGG19_model)
```



```
11/11 - 1s - loss: 1.3534 - accuracy: 0.8575 - 1s/epoch - 113ms/step
Restored model, accuracy: 85.75%
11/11 [=====] - 1s 57ms/step
```

Test Accuracy: 22.22%

Accuracy for class indian market is 83.95% consisting of 81 images  
Accuracy for class onion is 81.93% consisting of 83 images  
Accuracy for class potato is 72.84% consisting of 81 images  
Accuracy for class tomato is 100.0% consisting of 106 images



```
'''early_stopping_cb = tf.keras.callbacks.EarlyStopping(
    monitor = 'val_loss', patience = 5, restore_best_weights=True
)'''
pretrained_model=tf.keras.applications.ResNet101(weights='imagenet',in
clude_top=False,input_shape=(128,128,3))
pretrained_model.trainable=False
ResNet101_model=tf.keras.Sequential([
    pretrained_model,
    #layers.Flatten(),
    layers.GlobalAveragePooling2D(),
    layers.Dense(4,activation='softmax')
])
ResNet101_model.compile(optimizer='Adam',
metrics=['accuracy','Precision','Recall'],
loss='categorical_crossentropy')

history=ResNet101_model.fit(train_ds,epochs=20,validation_data=val_ds)
#pretrained_model.summary()

Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/resnet/
resnet101_weights_tf_dim_ordering_tf_kernels_notop.h5
171446536/171446536 [=====] - 9s 0us/step
Epoch 1/20
79/79 [=====] - 21s 158ms/step - loss: 0.4259
```

```
- accuracy: 0.8361 - precision: 0.8530 - recall: 0.8170 - val_loss:
0.2283 - val_accuracy: 0.9171 - val_precision: 0.9241 - val_recall:
0.9123
Epoch 2/20
79/79 [=====] - 9s 103ms/step - loss: 0.1343
- accuracy: 0.9498 - precision: 0.9550 - recall: 0.9482 - val_loss:
0.2086 - val_accuracy: 0.9314 - val_precision: 0.9343 - val_recall:
0.9298
Epoch 3/20
79/79 [=====] - 8s 96ms/step - loss: 0.0904 -
accuracy: 0.9701 - precision: 0.9735 - recall: 0.9677 - val_loss:
0.2161 - val_accuracy: 0.9314 - val_precision: 0.9373 - val_recall:
0.9298
Epoch 4/20
79/79 [=====] - 9s 97ms/step - loss: 0.0689 -
accuracy: 0.9829 - precision: 0.9832 - recall: 0.9777 - val_loss:
0.1729 - val_accuracy: 0.9410 - val_precision: 0.9469 - val_recall:
0.9394
Epoch 5/20
79/79 [=====] - 10s 115ms/step - loss: 0.0453
- accuracy: 0.9884 - precision: 0.9888 - recall: 0.9872 - val_loss:
0.1598 - val_accuracy: 0.9522 - val_precision: 0.9521 - val_recall:
0.9506
Epoch 6/20
79/79 [=====] - 9s 111ms/step - loss: 0.0371
- accuracy: 0.9916 - precision: 0.9920 - recall: 0.9912 - val_loss:
0.1935 - val_accuracy: 0.9426 - val_precision: 0.9501 - val_recall:
0.9410
Epoch 7/20
79/79 [=====] - 8s 97ms/step - loss: 0.0263 -
accuracy: 0.9956 - precision: 0.9956 - recall: 0.9952 - val_loss:
0.1586 - val_accuracy: 0.9537 - val_precision: 0.9551 - val_recall:
0.9506
Epoch 8/20
79/79 [=====] - 10s 105ms/step - loss: 0.0221
- accuracy: 0.9964 - precision: 0.9968 - recall: 0.9964 - val_loss:
0.1693 - val_accuracy: 0.9490 - val_precision: 0.9518 - val_recall:
0.9458
Epoch 9/20
79/79 [=====] - 10s 114ms/step - loss: 0.0169
- accuracy: 0.9984 - precision: 0.9984 - recall: 0.9980 - val_loss:
0.1559 - val_accuracy: 0.9569 - val_precision: 0.9584 - val_recall:
0.9553
Epoch 10/20
79/79 [=====] - 9s 108ms/step - loss: 0.0137
- accuracy: 0.9996 - precision: 0.9996 - recall: 0.9992 - val_loss:
0.1663 - val_accuracy: 0.9506 - val_precision: 0.9535 - val_recall:
0.9490
Epoch 11/20
```

```
79/79 [=====] - 9s 99ms/step - loss: 0.0115 -  
accuracy: 0.9996 - precision: 0.9996 - recall: 0.9992 - val_loss:  
0.1670 - val_accuracy: 0.9522 - val_precision: 0.9536 - val_recall:  
0.9506  
Epoch 12/20  
79/79 [=====] - 10s 115ms/step - loss: 0.0100  
- accuracy: 1.0000 - precision: 1.0000 - recall: 1.0000 - val_loss:  
0.1633 - val_accuracy: 0.9553 - val_precision: 0.9569 - val_recall:  
0.9553  
Epoch 13/20  
79/79 [=====] - 11s 128ms/step - loss: 0.0100  
- accuracy: 1.0000 - precision: 1.0000 - recall: 1.0000 - val_loss:  
0.1774 - val_accuracy: 0.9458 - val_precision: 0.9471 - val_recall:  
0.9426  
Epoch 14/20  
79/79 [=====] - 10s 113ms/step - loss: 0.0082  
- accuracy: 1.0000 - precision: 1.0000 - recall: 1.0000 - val_loss:  
0.1691 - val_accuracy: 0.9537 - val_precision: 0.9553 - val_recall:  
0.9537  
Epoch 15/20  
79/79 [=====] - 9s 110ms/step - loss: 0.0078  
- accuracy: 0.9996 - precision: 0.9996 - recall: 0.9996 - val_loss:  
0.1687 - val_accuracy: 0.9537 - val_precision: 0.9553 - val_recall:  
0.9537  
Epoch 16/20  
79/79 [=====] - 10s 109ms/step - loss: 0.0062  
- accuracy: 1.0000 - precision: 1.0000 - recall: 1.0000 - val_loss:  
0.1743 - val_accuracy: 0.9537 - val_precision: 0.9537 - val_recall:  
0.9537  
Epoch 17/20  
79/79 [=====] - 10s 115ms/step - loss: 0.0056  
- accuracy: 1.0000 - precision: 1.0000 - recall: 1.0000 - val_loss:  
0.1700 - val_accuracy: 0.9537 - val_precision: 0.9552 - val_recall:  
0.9522  
Epoch 18/20  
79/79 [=====] - 9s 112ms/step - loss: 0.0050  
- accuracy: 1.0000 - precision: 1.0000 - recall: 1.0000 - val_loss:  
0.1785 - val_accuracy: 0.9506 - val_precision: 0.9506 - val_recall:  
0.9506  
Epoch 19/20  
79/79 [=====] - 9s 100ms/step - loss: 0.0047  
- accuracy: 1.0000 - precision: 1.0000 - recall: 1.0000 - val_loss:  
0.1692 - val_accuracy: 0.9522 - val_precision: 0.9522 - val_recall:  
0.9522  
Epoch 20/20  
79/79 [=====] - 10s 117ms/step - loss: 0.0042  
- accuracy: 1.0000 - precision: 1.0000 - recall: 1.0000 - val_loss:  
0.1818 - val_accuracy: 0.9522 - val_precision: 0.9522 - val_recall:  
0.9522
```

```

loss, acc, precision, recall = ResNet101_model.evaluate(test_ds,
verbose=2)
print("Restored model, accuracy: {:.2f}%".format(100 *
acc), 'precision:', precision, 'recall:', recall)

11/11 - 2s - loss: 0.5180 - accuracy: 0.8803 - precision: 0.8851 -
recall: 0.8775 - 2s/epoch - 209ms/step
Restored model, accuracy: 88.03% precision: 0.8850574493408203 recall:
0.8774929046630859

VGG19_model.metrics_names
['loss', 'accuracy']

ResNet101_model.metrics_names
['loss', 'accuracy', 'precision', 'recall']

def plot_accuracy_ResNet(model_fit):
    #accuracy graph

    x = range(0, len(model_fit.history.history['accuracy']))
    y_train = [acc * 100 for acc in
model_fit.history.history['accuracy']]
    y_val = [acc * 100 for acc in
model_fit.history.history['val_accuracy']]

    plt.plot(x, y_train, label='Train', color='b')
    #annot_max(x, y_train, xytext=(0.7,0.9))
    plt.plot(x, y_val, label='Val', color='r')
    #annot_max(x, y_val, xytext=(0.8,0.7))
    plt.ylabel('Accuracy', fontsize=15)
    plt.xlabel('epoch', fontsize=15)
    plt.legend()
    plt.show()

#plot_accuracy(ResNet101_model)
loss, accuracy, precision, recall = ResNet101_model.evaluate(test_ds,
verbose=2)
print("Restored model, accuracy: {:.2f}%".format(100 *
acc), 'precision:', precision, 'recall:', recall)

y_pred = ResNet101_model.predict(test_ds)
predicted_categories = tf.argmax(y_pred, axis=1)
true_cat = tf.concat([y for x, y in test_ds], axis=0)
true_categories = tf.argmax(true_cat, axis=1)

# calculate accuracy
test_acc = metrics.accuracy_score(true_categories,
predicted_categories) * 100
print(f'\nTest Accuracy: {test_acc:.2f}%\n')

```

```

classwise_accuracy(noise_path, 'indian market', ResNet101_model)
classwise_accuracy(onion_path, 'onion', ResNet101_model)
classwise_accuracy(potato_path, 'potato', ResNet101_model)
classwise_accuracy(tomato_path, 'tomato', ResNet101_model)

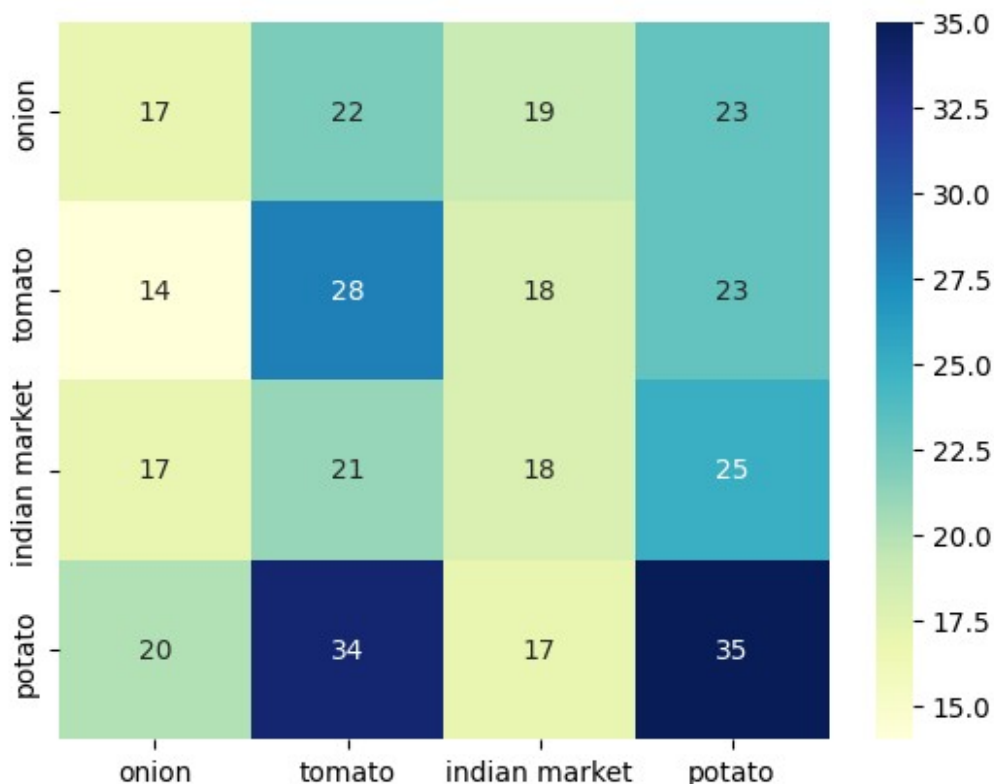
ConfusionMatrix(cnn_model, test_ds, class_dirs)

11/11 - 1s - loss: 0.5180 - accuracy: 0.8803 - precision: 0.8851 -
recall: 0.8775 - 1s/epoch - 122ms/step
Restored model, accuracy: 88.03% precision: 0.8850574493408203 recall:
0.8774929046630859
11/11 [=====] - 3s 67ms/step

Test Accuracy: 27.92%

```

Accuracy for class indian market is 81.48% consisting of 81 images  
 Accuracy for class onion is 90.36% consisting of 83 images  
 Accuracy for class potato is 76.54% consisting of 81 images  
 Accuracy for class tomato is 100.0% consisting of 106 images





# Testing our best model (ResNet)

ResNet achieved the highest accuracy of 94% among all the models we trained

Let's test our Finetuned ResNet to predict on some random unseen data to visualize how accurate it is !

```
ResNet101_model.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
resnet101 (Functional)	(None, 4, 4, 2048)	42658176
global_average_pooling2d_4 (GlobalAveragePooling2D)	(None, 2048)	0
dense_4 (Dense)	(None, 4)	8196

```
=====  
Total params: 42666372 (162.76 MB)  
Trainable params: 8196 (32.02 KB)  
Non-trainable params: 42658176 (162.73 MB)
```

```
grid_test_model(ResNet101_model)
```

```
1/1 [=====] - 0s 28ms/step
```

```
<ipython-input-21-40cb159e3305>:32: MatplotlibDeprecationWarning:  
Auto-removal of overlapping axes is deprecated since 3.6 and will be  
removed two minor releases later; explicitly call ax.remove() as  
needed.
```

```
plt.subplot(2, 4, n)
```

```
1/1 [=====] - 0s 179ms/step
```

```
1/1 [=====] - 0s 47ms/step
```

```
1/1 [=====] - 0s 37ms/step
```

prediction : potato prediction : tomato prediction : tomato prediction : tomato

0.0 % indian market	7.0 % indian market	0.0 % indian market	0.0 % indian market
0.0 % onion	0.0 % onion	0.0 % onion	0.0 % onion
100.0 % potato	1.0 % potato	0.0 % potato	0.0 % potato
0.0 % tomato	92.0 % tomato	100.0 % tomato	100.0 % tomato

